


NSBI signal-strength bias vs training statistics: fine-tuned pretrained EveNet vs scratch
48 weighted toys, 300 fb⁻¹



error bars: spread over seeds (3 seeds) or statistical error of the 48-toy mean (1 seed). Hollow marker: some fits failed; the point at -2.5: all 48 fits collapsed to $\hat{\mu} = 0$.